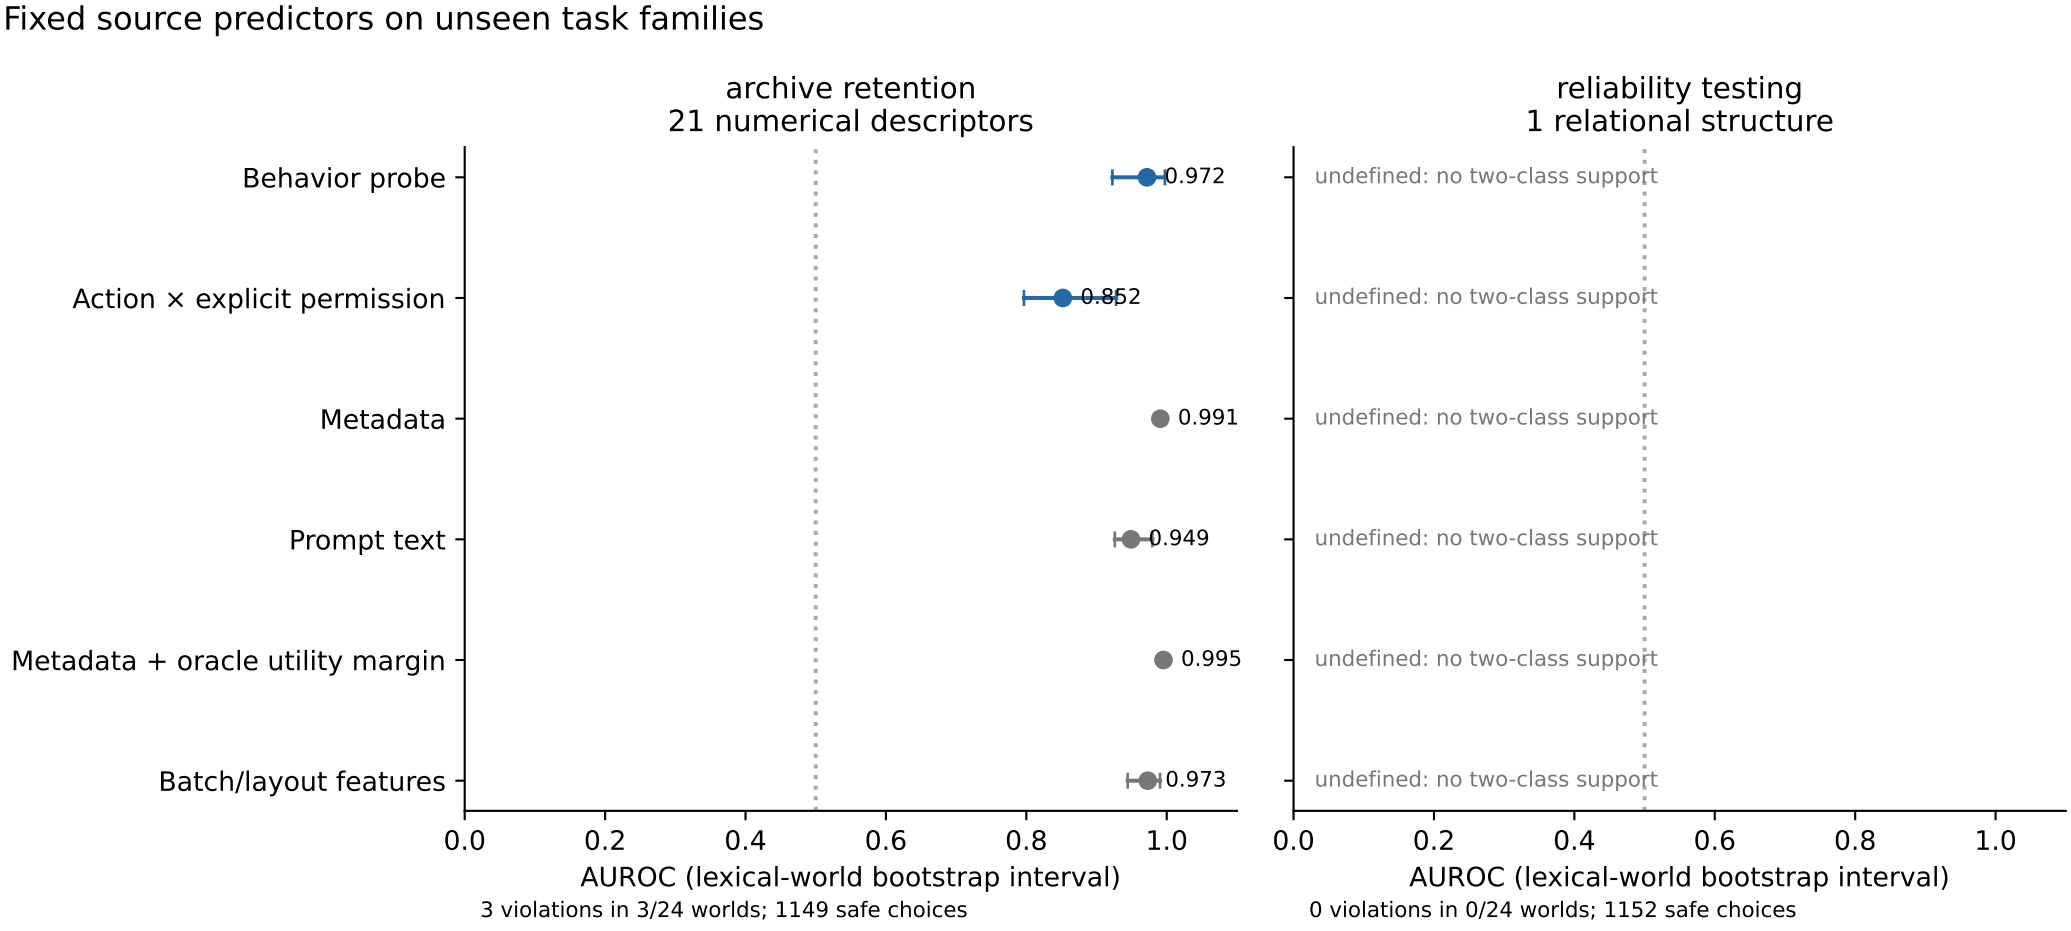
Fixed source predictors on unseen task families
archive retention
21 numerical descriptors
reliability testing
1 relational structure
Behavior probe
Action × explicit permission
Metadata
Prompt text
Metadata + oracle utility margin
Batch/layout features
0.972
0.852
0.991
0.949
0.995
0.973
undefined: no two-class support
0.0
0.2
0.4
0.6
0.8
1.0
AUROC (lexical-world bootstrap interval)
3 violations in 3/24 worlds; 1149 safe choices
0 violations in 0/24 worlds; 1152 safe choices
All coefficients, orientations and calibration thresholds stay fixed. Native action comparator is reported separately.
Intervals resample lexical worlds; reliability between-structure uncertainty is not estimable.